A Proposal on Standard Notation for Machine Learning

January 27, 2020

Summary

The field of machine learning is evolving rapidly in recent years. Communication between different researchers and research groups becomes increasingly important. A key challenge for communication arises from inconsistent notation usages among different papers. This proposal suggests a standard for commonly used mathematical notation for machine learning. In this first version, only some notation are mentioned and more notation are left to be done. This proposal will be regularly updated based on the progress of the field. We look forward to more suggestions to improve this proposal in future versions.

Contents

1 Dataset

Dataset $S = \{\boldsymbol{z}_i\}_{i=1}^n = \{(\boldsymbol{x}_i, \boldsymbol{y}_i)\}_{i=1}^n$ is sampled from a distribution $\mathcal{D}$ over a domain $\mathcal{Z} = \mathcal{X} \times \mathcal{Y}$.

$\mathcal{X}$ is the instance domain (a set), $\mathcal{Y}$ is the label domain (a set), and $\mathcal{Z} = \mathcal{X} \times \mathcal{Y}$ is the example domain (a set).

Usually, $\mathcal{X}$ is a subset of $\mathbb{R}^d$ and $\mathcal{Y}$ is a subset of $\mathbb{R}^{d_o}$, where d is the input dimension, d_o is the output dimension.

$n = \#S$ is the number of samples. Without specification, S and n are for the training set.

2 Function

Hypothesis space is denoted by $\mathcal{H}$. Hypothesis function is denoted by $f_{\boldsymbol{\theta}}(\boldsymbol{x}) \in \mathcal{H}$ or $f(\boldsymbol{x}; \boldsymbol{\theta}) \in \mathcal{H}$ with $f_{\boldsymbol{\theta}} : \mathcal{X} \to \mathcal{Y}$.

$\boldsymbol{\theta}$ denotes the set of parameters of $f_{\boldsymbol{\theta}}$.

If there exists a target function, it is denoted by f^* or $f : \mathcal{X} \to \mathcal{Y}$ satisfying $\boldsymbol{y}_i = f^*(\boldsymbol{x}_i)$ for $i = 1, \ldots, n$.

3 Loss function

Loss function, denoted by $\ell : \mathcal{H} \times \mathcal{Z} \to \mathbb{R}_+ := [0, +\infty)$, measures the difference between a predicted label and a true label, e.g., L^2 loss:

$$\ell(f_{\boldsymbol{\theta}}, \boldsymbol{z}) = \frac{1}{2}(f_{\boldsymbol{\theta}}(\boldsymbol{x}) - \boldsymbol{y})^2,$$

where $\boldsymbol{z} = (\boldsymbol{x}, \boldsymbol{y})$. $\ell(f_{\boldsymbol{\theta}}, \boldsymbol{z})$ can also be written as

$$\ell(f_{\boldsymbol{\theta}}(\boldsymbol{x}), \boldsymbol{y})$$

for convenience. Empirical risk or training loss for a set $S = \{(\boldsymbol{x}_i, \boldsymbol{y}_i)\}_{i=1}^n$ is denoted by $L_S(\boldsymbol{\theta})$ or $L_n(\boldsymbol{\theta})$ or $\mathcal{R}_n(\boldsymbol{\theta})$ or $\mathcal{R}_S(\boldsymbol{\theta})$,

$$L_S(\boldsymbol{\theta}) = \frac{1}{n}\sum_{i=1}^n \ell(f_{\boldsymbol{\theta}}(\boldsymbol{x}_i), \boldsymbol{y}_i). \tag{1}$$

Without ambiguity, L is also used for L_S.

The population risk or expected loss is denoted by

$$L_{\mathcal{D}}(\boldsymbol{\theta}) = \mathbb{E}_{\mathcal{D}}\ell(f_{\boldsymbol{\theta}}(\boldsymbol{x}), \boldsymbol{y}), \tag{2}$$

where $\boldsymbol{z} = (\boldsymbol{x}, \boldsymbol{y})$ follows the distribution $\mathcal{D}$.

4 Activation function

Activation function is denoted by $\sigma(x)$.

Example 1. *Some commonly used activation functions are*

    1. $\sigma(x) = \mathrm{ReLU}(x) = \max(0, x)$;

    2. $\sigma(x) = \mathrm{sigmoid}(x) = \frac{1}{1+\mathrm{e}^{-x}}$;

    3. $\sigma(x) = \tanh x$;

    4. $\sigma(x) = \cos x, \sin x$.

5 Two-layer neural network

The neuron number of the hidden layer is denoted by m. The two-layer neural network is

$$f_{\boldsymbol{\theta}}(\boldsymbol{x}) = \sum_{j=1}^{m} a_j \sigma(\boldsymbol{w}_j \cdot \boldsymbol{x} + b_j), \tag{3}$$

where σ is the activation function, $\boldsymbol{w}_j$ is the input weight, a_j is the output weight, b_j is the bias term. We denote the set of parameters by

$$\boldsymbol{\theta} = (a_1, \cdots, a_m, \boldsymbol{w}_1, \cdots, \boldsymbol{w}_m, b_1, \cdots, b_m).$$

6 General deep neural network

The counting of the layer number excludes the input layer. A L-layer neural network is denoted by

$$f_{\boldsymbol{\theta}}(\boldsymbol{x}) = \boldsymbol{W}^{[L-1]}\sigma \circ (\boldsymbol{W}^{[L-2]}\sigma \circ (\cdots (\boldsymbol{W}^{[1]}\sigma \circ (\boldsymbol{W}^{[0]}\boldsymbol{x} + \boldsymbol{b}^{[0]}) + \boldsymbol{b}^{[1]}) \cdots) + \boldsymbol{b}^{[L-2]}) + \boldsymbol{b}^{[L-1]}, \tag{4}$$

where $\boldsymbol{W}^{[l]} \in \mathbb{R}^{m_{l+1} \times m_l}$, $\boldsymbol{b}^{[l]} = \mathbb{R}^{m_{l+1}}$, $m_0 = d_{\mathrm{in}} = d$, $m_L = d_{\mathrm{o}}$, σ is a scalar function and "∘" means entry-wise operation. We denote the set of parameters by

$$\boldsymbol{\theta} = (\boldsymbol{W}^{[0]}, \boldsymbol{W}^{[1]}, \ldots, \boldsymbol{W}^{[L-1]}, \boldsymbol{b}^{[0]}, \boldsymbol{b}^{[1]}, \ldots, \boldsymbol{b}^{[L-1]}),$$

and an entry of $\boldsymbol{W}^{[l]}$ by $\boldsymbol{W}_{ij}^{[l]}$. This can also be defined recursively.

$$f_{\boldsymbol{\theta}}^{[0]}(\boldsymbol{x}) = \boldsymbol{x}, \tag{5}$$

$$f_{\boldsymbol{\theta}}^{[l]}(\boldsymbol{x}) = \sigma \circ (\boldsymbol{W}^{[l-1]}f_{\boldsymbol{\theta}}^{[l-1]}(\boldsymbol{x}) + \boldsymbol{b}^{[l-1]}) \quad 1 \le l \le L-1, \tag{6}$$

$$f_{\boldsymbol{\theta}}(\boldsymbol{x}) = f_{\boldsymbol{\theta}}^{[L]}(\boldsymbol{x}) = \boldsymbol{W}^{[L-1]}f_{\boldsymbol{\theta}}^{[L-1]}(\boldsymbol{x}) + \boldsymbol{b}^{[L-1]}. \tag{7}$$

7 Complexity

The VC-dimension of a hypothesis class $\mathcal{H}$ is denoted $\text{VCdim}(\mathcal{H})$.

The Rademacher complexity of a hypothesis space $\mathcal{H}$ on a sample set S is denoted by $\text{Rad}(\mathcal{H} \circ S)$ or $\text{Rad}_S(\mathcal{H})$. The complexity $\text{Rad}_S(\mathcal{H})$ is random because of the randomness of S. The expectation of the empirical Rademacher complexity over all samples of size n is denoted by

$$\text{Rad}_n(\mathcal{H}) = \mathbb{E}_S \text{Rad}_S(\mathcal{H}).$$

8 Training

The Gradient Descent is often denoted by GD. The Stochastic Gradient Descent is often denoted by SGD.

A batch set is denoted by B and the batch size is denoted by b.

The learning rate is denoted by η.

9 Fourier Frequency

The discretized frequency is denoted by $\boldsymbol{k}$, and the continuous frequency is denoted by $\boldsymbol{\xi}$.

10 Convolution

The convolution operation is denoted by $*$.

11 Notation table

symbol	meaning	LaTeX	simplied
$\boldsymbol{x}$	input	\bm{x}	\vx
$\boldsymbol{y}$	output, label	\bm{y}	\vy
d	input dimension	d	
$d_{\rm o}$	output dimension	d_{\rm o}	
n	number of samples	n	
$\mathcal{X}$	instances domain (a set)	\mathcal{X}	\fX
$\mathcal{Y}$	labels domain (a set)	\mathcal{Y}	\fY
$\mathcal{Z}$	$= \mathcal{X} \times \mathcal{Y}$ example domain	\mathcal{Z}	\fZ
$\mathcal{H}$	hypothesis space (a set)	\mathcal{H}	\fH
$\boldsymbol{\theta}$	a set of parameters	\bm{\theta}	\vtheta
$f_{\boldsymbol{\theta}} : \mathcal{X} \to \mathcal{Y}$	hypothesis function	\f_{\bm{\theta}}	f_{\vtheta}
f or $f^* : \mathcal{X} \to \mathcal{Y}$	target function	f,f^*	
$\ell : \mathcal{H} \times \mathcal{Z} \to \mathbb{R}^+$	loss function	\ell	
$\mathcal{D}$	distribution of $\mathcal{Z}$	\mathcal{D}	\fD
$S = \{\boldsymbol{z}_i\}_{i=1}^n$	$= \{(\boldsymbol{x}_i, \boldsymbol{y}_i)\}_{i=1}^n$ sample set		
$L_S(\boldsymbol{\theta})$, $L_n(\boldsymbol{\theta})$, $\mathcal{R}_n(\boldsymbol{\theta})$, $\mathcal{R}_S(\boldsymbol{\theta})$	empirical risk or training loss		
$L_{\mathcal{D}}(\boldsymbol{\theta})$	population risk or expected loss		
$\sigma : \mathbb{R} \to \mathbb{R}^+$	activation function	\sigma	
$\boldsymbol{w}_j$	input weight	\bm{w}_j	\vw_j
a_j	output weight	a_j	
b_j	bias term	b_j	
$f_{\boldsymbol{\theta}}(\boldsymbol{x})$ or $f(\boldsymbol{x}; \boldsymbol{\theta})$	neural network	f_{\bm{\theta}}	f_{\vtheta}
$\sum_{j=1}^m a_j \sigma(\boldsymbol{w}_j \cdot \boldsymbol{x} + b_j)$	two-layer neural network		
$\mathrm{VCdim}(\mathcal{H})$	VC-dimension of $\mathcal{H}$		
$\mathrm{Rad}(\mathcal{H} \circ S)$, $\mathrm{Rad}_S(\mathcal{H})$	Rademacher complexity of $\mathcal{H}$ on S		
$\mathrm{Rad}_n(\mathcal{H})$	Rademacher complexity over samples of size n		
GD	gradient descent		
SGD	stochastic gradient descent		
B	a batch set	B	
b	batch size	b	
η	learning rate	\eta	
$\boldsymbol{k}$	discretized frequency	\bm{k}	\vk
$\boldsymbol{\xi}$	continuous frequency	\bm{\xi}	\vxi
$*$	convolution operation	*	

12 L-layer neural network

symbol	meaning	LaTeX	simplied
d	input dimension	`d`	
$d_{\rm o}$	output dimension	`d_{\rm o}`	
m_l	the number of lth layer neuron, $m_0 = d$, $m_L = d_{\rm o}$	`m_l`	
$\boldsymbol{W}^{[l]}$	the lth layer weight	`\bm{W}^{[l]}`	`\mW^{[l]}`
$\boldsymbol{b}^{[l]}$	the lth layer bias term	`\bm{b}^{[l]}`	`\vb^{[l]}`
$\circ$	entry-wise operation	`\circ`	
$\sigma : \mathbb{R} \to \mathbb{R}^+$	activation function	`\sigma`	
$\boldsymbol{\theta}$	$= (\boldsymbol{W}^{[0]}, \ldots, \boldsymbol{W}^{[L-1]}, \boldsymbol{b}^{[0]}, \ldots, \boldsymbol{b}^{[L-1]})$, parameters	`\bm{\theta}`	`\vtheta`
$f_{\boldsymbol{\theta}}^{[0]}(\boldsymbol{x})$	$= \boldsymbol{x}$		
$f_{\boldsymbol{\theta}}^{[l]}(\boldsymbol{x})$	$= \sigma \circ (\boldsymbol{W}^{[l-1]} f_{\boldsymbol{\theta}}^{[l-1]}(\boldsymbol{x}) + \boldsymbol{b}^{[l-1]})$, l-th layer output		
$f_{\boldsymbol{\theta}}(\boldsymbol{x})$	$= f_{\boldsymbol{\theta}}^{[L]}(\boldsymbol{x}) = \boldsymbol{W}^{[L-1]} f_{\boldsymbol{\theta}}^{[L-1]}(\boldsymbol{x}) + \boldsymbol{b}^{[L-1]}$, L-layer NN		

13   Acknowledgements